



Numpy

Peerapon Vateekul, Ph.D.

peerapon.v@4amconsult.com



Outlines

- **Numpy**
- **Numpy vs Python List**
- **Numpy Array**
- **Numpy Array Indexing**
- **Numpy Math Operations**
- **Numpy Operation Summary**
 - Creating Numpy
 - Inspecting Properties
 - Copying/sorting/reshaping
 - Working with scalar
 - Working with 2 vectors
 - Advanced Math Functions
 - Statistics
- **Broadcasting**



Outlines

- **Numpy**
- **Numpy vs Python List**
- **Numpy Array**
- **Numpy Array Indexing**
- **Numpy Math Operations**
- **Numpy Operation Summary**
 - Creating Numpy
 - Inspecting Properties
 - Copying/sorting/reshaping
 - Working with scalar
 - Working with 2 vectors
 - Advanced Math Functions
 - Statistics
- **Broadcasting**



Numpy

- Numpy is the core library for scientific computing in Python.
- It provides a high-performance multidimensional array object, and tools for working with these arrays.
- If you are already familiar with MATLAB, you might find this tutorial useful to get started with Numpy.

+ Numpy vs Python List

- **Size** - Numpy data structures take up **less space**
- **Performance** - they have a need for speed and are **faster** than lists
- **Functionality** - SciPy and NumPy have **optimized functions** such as linear algebra operations built in.

```
import time
import numpy as np

size_of_vec = 1000000 #1M

def pure_python():
    t1 = time.time()
    x = range(size_of_vec)
    y = range(size_of_vec)
    z = [x[i] + y[i] for i in range (len(x))]
    t2 = time.time()
    used_time = t2-t1
    print(used_time, 'seconds')
    return used_time

def numpy_version():
    t1 = time.time()
    x = np.arange(size_of_vec)
    y = np.arange(size_of_vec)
    z = x+y
    t2 = time.time()
    used_time = t2-t1
    print(used_time, 'seconds')
    return used_time

t1 = pure_python()
t2 = numpy_version()
print('faster', t1 - t2, 'seconds')
print('faster', t1/t2, 'times')
```

```
0.3190653324127197 seconds
0.0070002079010009766 seconds
faster 0.31206512451171875 seconds
faster 45.579408058308644 times
```



Numpy vs Python List


- Python list does not support matrix operations

```
a = [1,3,5,7,9]
b = [3,5,6,7,9]
c = a + b
print( c )
[1, 3, 5, 7, 9, 3, 5, 6, 7, 9]
```



```
def add_vector(a, b):
    c = [ a[i]+b[i] for i in range(len(a)) ]
    return c
```

```
a = [1,3,5,7,9]
b = [3,5,6,7,9]
c = add_vector(a,b)
print( c )
```



```
import numpy as np
```

```
a = np.array([1,3,5,7,9])
b = np.array([3,5,6,7,9])
c = a + b
```



Outlines

- **Numpy**
- **Numpy vs Python List**
- **Numpy Array**
- **Numpy Array Indexing**
- **Numpy Math Operations**
- **Numpy Operation Summary**
 - Creating Numpy
 - Inspecting Properties
 - Copying/sorting/reshaping
 - Working with scalar
 - Working with 2 vectors
 - Advanced Math Functions
 - Statistics
- **Broadcasting**



Numpy Array

- A numpy array is a grid of values, all of the same type, and is indexed by a tuple of nonnegative integers. The number of dimensions is the *rank* of the array; the *shape* of an array is a tuple of integers giving the size of the array along each dimension.
- We can initialize numpy arrays from nested Python lists, and access elements using square brackets:

```
import numpy as np

a = np.array([1, 2, 3])    # Create a rank 1 array
print(type(a))            # Prints "<class 'numpy.ndarray'>"
print(a.shape)            # Prints "(3,)"
print(a[0], a[1], a[2])   # Prints "1 2 3"
a[0] = 5                  # Change an element of the array
print(a)                  # Prints "[5, 2, 3]"

b = np.array([[1,2,3],[4,5,6]]) # Create a rank 2 array
print(b.shape)            # Prints "(2, 3)"
print(b[0, 0], b[0, 1], b[1, 0]) # Prints "1 2 4"
```


Numpy Array

- Numpy also provides many functions to create arrays:

```
import numpy as np

a = np.zeros((2,2))    # Create an array of all zeros
print(a)              # Prints "[[ 0.  0.]
                      #           [ 0.  0.]]"

b = np.ones((1,2))    # Create an array of all ones
print(b)              # Prints "[[ 1.  1.]]"

c = np.full((2,2), 7)  # Create a constant array
print(c)              # Prints "[[ 7.  7.]
                      #           [ 7.  7.]]"

d = np.eye(2)         # Create a 2x2 identity matrix
print(d)              # Prints "[[ 1.  0.]
                      #           [ 0.  1.]]"

e = np.random.random((2,2))  # Create an array filled with random values
print(e)                # Might print "[[ 0.91940167  0.08143941]
                      #           [ 0.68744134  0.87236687]]"
```

+ Numpy Array Indexing

- **1) Integer array indexing:** When you index into numpy arrays using slicing, the resulting array view will always be a subarray of the original array. In contrast, integer array indexing allows you to construct arbitrary arrays using the data from another array. Here is an example:

```
import numpy as np

a = np.array([[1,2], [3, 4], [5, 6]])

# An example of integer array indexing.
# The returned array will have shape (3,) and
print(a[[0, 1, 2], [0, 1, 0]]) # Prints "[1 4 5]"

# The above example of integer array indexing is equivalent to this:
print(np.array([a[0, 0], a[1, 1], a[2, 0]])) # Prints "[1 4 5]"

# When using integer array indexing, you can reuse the same
# element from the source array:
print(a[[0, 0], [1, 1]]) # Prints "[2 2]"

# Equivalent to the previous integer array indexing example
print(np.array([a[0, 1], a[0, 1]])) # Prints "[2 2]"
```

+ Numpy Array Indexing

- Numpy offers several ways to index into arrays.
- **2) Slicing:** Similar to Python lists, numpy arrays can be sliced. Since arrays may be multidimensional, you must specify a slice for each dimension of the array:

```
import numpy as np

# Create the following rank 2 array with shape (3, 4)
# [[ 1  2  3  4]
#  [ 5  6  7  8]
#  [ 9 10 11 12]]
a = np.array([[1,2,3,4], [5,6,7,8], [9,10,11,12]])

# Use slicing to pull out the subarray consisting of the first 2 rows
# and columns 1 and 2; b is the following array of shape (2, 2):
# [[2 3]
#  [6 7]]
b = a[:2, 1:3]

# A slice of an array is a view into the same data, so modifying it
# will modify the original array.
print(a[0, 1])    # Prints "2"
b[0, 0] = 77      # b[0, 0] is the same piece of data as a[0, 1]
print(a[0, 1])    # Prints "77"
```

+ Numpy Array Indexing

- You can also **mix integer indexing with slice indexing**. However, doing so will yield an array of lower rank than the original array. Note that this is quite different from the way that MATLAB handles array slicing:

```
import numpy as np

# Create the following rank 2 array with shape (3, 4)
# [[ 1  2  3  4]
#  [ 5  6  7  8]
#  [ 9 10 11 12]]
a = np.array([[1,2,3,4], [5,6,7,8], [9,10,11,12]])

# Two ways of accessing the data in the middle row of the array.
# Mixing integer indexing with slices yields an array of lower rank,
# while using only slices yields an array of the same rank as the
# original array:
row_r1 = a[1, :]    # Rank 1 view of the second row of a
row_r2 = a[1:2, :]  # Rank 2 view of the second row of a
print(row_r1, row_r1.shape)  # Prints "[5 6 7 8] (4,)"
print(row_r2, row_r2.shape)  # Prints "[[5 6 7 8]] (1, 4)"

# We can make the same distinction when accessing columns of an array:
col_r1 = a[:, 1]
col_r2 = a[:, 1:2]
print(col_r1, col_r1.shape)  # Prints "[ 2  6 10] (3,)"
print(col_r2, col_r2.shape)  # Prints "[[ 2]
                             #          [ 6]
                             #          [10]] (3, 1)"
```

+ Numpy Array Indexing

- One useful trick with integer array indexing is selecting or mutating one element from each row of a matrix:

```
import numpy as np

# Create a new array from which we will select elements
a = np.array([[1,2,3], [4,5,6], [7,8,9], [10, 11, 12]])

print(a) # prints "array([[ 1,  2,  3],
#           [ 4,  5,  6],
#           [ 7,  8,  9],
#           [10, 11, 12]])"
```

```
# Create an array of indices
b = np.array([0, 2, 0, 1])

# Select one element from each row of a using the indices in b
print(a[np.arange(4), b]) # Prints "[ 1  6  7 11]"

# Mutate one element from each row of a using the indices in b
a[np.arange(4), b] += 10

print(a) # prints "array([[11,  2,  3],
#           [ 4,  5, 16],
#           [17,  8,  9],
#           [10, 21, 12]])"
```


+ Numpy Array Indexing

- **3) Boolean array indexing:** Boolean array indexing lets you pick out arbitrary elements of an array. Frequently this type of indexing is used to select the elements of an array that satisfy some condition. Here is an example:

```
import numpy as np

a = np.array([[1,2], [3, 4], [5, 6]])

bool_idx = (a > 2)    # Find the elements of a that are bigger than 2;
                      # this returns a numpy array of Booleans of the same
                      # shape as a, where each slot of bool_idx tells
                      # whether that element of a is > 2.

print(bool_idx)        # Prints "[[False False]
                      #      [ True  True]
                      #      [ True  True]]"

# We use boolean array indexing to construct a rank 1 array
# consisting of the elements of a corresponding to the True values
# of bool_idx
print(a[bool_idx])     # Prints "[3 4 5 6]"

# We can do all of the above in a single concise statement:
print(a[a > 2])        # Prints "[3 4 5 6]"
```



Outlines

- **Numpy**
- **Numpy vs Python List**
- **Numpy Array**
- **Numpy Array Indexing**
- **Numpy Math Operations**
- **Numpy Operation Summary**
 - Creating Numpy
 - Inspecting Properties
 - Copying/sorting/reshaping
 - Working with scalar
 - Working with 2 vectors
 - Advanced Math Functions
 - Statistics
- **Broadcasting**

+ Numpy Math Operations

- Basic mathematical functions operate elementwise on arrays, and are available both as operator overloads and as functions in the numpy module:

```
import numpy as np

x = np.array([[1,2],[3,4]], dtype=np.float64)
y = np.array([[5,6],[7,8]], dtype=np.float64)

# Elementwise sum; both produce the array
# [[ 6.0  8.0]
#  [10.0 12.0]]
print(x + y)
print(np.add(x, y))

# Elementwise difference; both produce the array
# [[-4.0 -4.0]
#  [-4.0 -4.0]]
print(x - y)
print(np.subtract(x, y))
```

```
# Elementwise product; both produce the array
# [[ 5.0 12.0]
#  [21.0 32.0]]
print(x * y)
print(np.multiply(x, y))

# Elementwise division; both produce the array
# [[ 0.2          0.33333333]
#  [ 0.42857143  0.5         ]]
print(x / y)
print(np.divide(x, y))

# Elementwise square root; produces the array
# [[ 1.          1.41421356]
#  [ 1.73205081  2.         ]]
print(np.sqrt(x))
```


+ Numpy Math Operations

- Note that unlike MATLAB, `*` is elementwise multiplication, not matrix multiplication. We instead use the `dot` function to compute inner products of vectors, to multiply a vector by a matrix, and to multiply matrices. `dot` is available both as a function in the `numpy` module and as an instance method of array objects:

```
import numpy as np

x = np.array([[1,2],[3,4]])
y = np.array([[5,6],[7,8]])

v = np.array([9,10])
w = np.array([11, 12])

# Inner product of vectors; both produce 219
print(v.dot(w))
print(np.dot(v, w))
```

```
# Matrix / vector product; both produce the rank 1 array [29 67]
print(x.dot(v))
print(np.dot(x, v))

# Matrix / matrix product; both produce the rank 2 array
# [[19 22]
#  [43 50]]
print(x.dot(y))
print(np.dot(x, y))
```

+ Numpy Math Operations

- Numpy provides many useful functions for performing computations on arrays; one of the most useful is sum:

```
import numpy as np

x = np.array([[1,2],[3,4]])

print(np.sum(x))    # Compute sum of all elements; prints "10"
print(np.sum(x, axis=0)) # Compute sum of each column; prints "[4 6]"
print(np.sum(x, axis=1)) # Compute sum of each row; prints "[3 7]"
```

+ Numpy Math Operations

- Apart from computing mathematical functions using arrays, we frequently need to reshape or otherwise manipulate data in arrays. The simplest example of this type of operation is transposing a matrix; to transpose a matrix, simply use the **T** attribute of an array object:

```
import numpy as np

x = np.array([[1,2], [3,4]])
print(x)      # Prints "[[1 2]
               #           [3 4]]"
print(x.T)    # Prints "[[1 3]
               #           [2 4]]"

# Note that taking the transpose of a rank 1 array does nothing:
v = np.array([1,2,3])
print(v)      # Prints "[1 2 3]"
print(v.T)    # Prints "[1 2 3]"
```



Outlines

- **Numpy**
- **Numpy vs Python List**
- **Numpy Array**
- **Numpy Array Indexing**
- **Numpy Math Operations**
- **Numpy Operation Summary**
 - Creating Numpy
 - Inspecting Properties
 - Copying/sorting/reshaping
 - Working with scalar
 - Working with 2 vectors
 - Advanced Math Functions
 - Statistics
- **Broadcasting**

+ Numpy Operation Summary

- Creating Numpy

- `np.array([1,2,3])` | One dimensional array
- `np.array([(1,2,3),(4,5,6)])` | Two dimensional array
- `np.zeros(3)` | 1D array of length 3 all values 0
- `np.ones(3)` | 3 array with all values 1
- `np.linspace(0,100,6)` | Array of 6 evenly divided values from 0 to 100
- `np.arange(0,10,3)` | Array of values from 0 to less than 10 with step 3 (eg [0,3,6,9])
- `np.random.rand(4,5)` | 4x5 array of random floats between 0–1
- `np.random.randint(5, size=(3))` | 3 array with random ints between 0–4

+ Numpy Operation Summary

- Inspecting Properties

- `arr.size` | Returns number of elements in arr
- `arr.shape` | Returns dimensions of arr (rows,columns)
- `arr.dtype` | Returns type of elements in arr
- `arr.astype(dtype)` | Convert arr elements to type dtype
- `arr.tolist()` | Convert arr to a Python list



Numpy Operation Summary

- Copying/Sorting/Reshaping

- `np.copy(arr)` | Copies arr to new memory
- `arr.view(dtype)` | Creates view of arr elements with type dtype
- `arr.sort()` | Sorts arr
- `arr.sort(axis=0)` | Sorts specific axis of arr
- `two_d_arr.flatten()` | Flattens 2D array two_d_arr to 1D
- `arr.T` | Transposes arr (rows become columns and vice versa)
- `arr.reshape(3,4)` | Reshapes arr to 3 rows, 4 columns without changing data
- `arr.resize((5,6))` | Changes arr shape to 5x6 and fills new values with 0

+ Numpy Operation Summary

- Working with Scalar

- `np.add(arr, 1)` | Add 1 to each array element
- `np.subtract(arr, 2)` | Subtract 2 from each array element
- `np.multiply(arr, 3)` | Multiply each array element by 3
- `np.divide(arr, 4)` | Divide each array element by 4 (returns `np.nan` for division by zero)
- `np.power(arr, 5)` | Raise each array element to the 5th power
- `np.sqrt(arr)` | Square root of each element in the array

+ Numpy Operation Summary

- Working with 2 Vectors

- `np.add(arr1,arr2)` | Elementwise add arr2 to arr1
- `np.subtract(arr1,arr2)` | Elementwise subtract arr2 from arr1
- `np.multiply(arr1,arr2)` | Elementwise multiply arr1 by arr2
- `np.divide(arr1,arr2)` | Elementwise divide arr1 by arr2
- `np.power(arr1,arr2)` | Elementwise raise arr1 raised to the power of arr2
- `np.array_equal(arr1,arr2)` | Returns True if the arrays have the same elements and shape

+ Numpy Operation Summary

- Advanced Math Functions

- `np.sin(arr)` | Sine of each element in the array
- `np.log(arr)` | Natural log of each element in the array
- `np.abs(arr)` | Absolute value of each element in the array
- `np.ceil(arr)` | Rounds up to the nearest int
- `np.floor(arr)` | Rounds down to the nearest int
- `np.round(arr)` | Rounds to the nearest int

+ Numpy Operation Summary

- Statistics

- `np.mean(arr,axis=0)` | Returns mean along specific axis
- `arr.sum()` | Returns sum of arr
- `arr.min()` | Returns minimum value of arr
- `arr.max(axis=0)` | Returns maximum value of specific axis
- `np.var(arr)` | Returns the variance of array
- `np.std(arr,axis=1)` | Returns the standard deviation of specific axis
- `arr.corrcoef()` | Returns correlation coefficient of array



Outlines

- **Numpy**
- **Numpy vs Python List**
- **Numpy Array**
- **Numpy Array Indexing**
- **Numpy Math Operations**
- **Numpy Operation Summary**
 - Creating Numpy
 - Inspecting Properties
 - Copying/sorting/reshaping
 - Working with scalar
 - Working with 2 vectors
 - Advanced Math Functions
 - Statistics
- **Broadcasting**

Broadcasting

- Broadcasting is a powerful mechanism that allows numpy to work with arrays of different shapes when performing arithmetic operations. Frequently we have a smaller array and a larger array, and we want to use the smaller array multiple times to perform some operation on the larger array.
- For example, suppose that we want to add a constant vector to each row of a matrix. We could do it like this:

```
import numpy as np

# We will add the vector v to each row of the matrix x,
# storing the result in the matrix y
x = np.array([[1,2,3], [4,5,6], [7,8,9], [10, 11, 12]])
v = np.array([1, 0, 1])
y = np.empty_like(x)    # Create an empty matrix with the same shape as x
```

Broadcasting

- Broadcasting is a powerful mechanism that allows numpy to work with arrays of different shapes when performing arithmetic operations. Frequently we have a smaller array and a larger array, and we want to use the smaller array multiple times to perform some operation on the larger array.
- For example, suppose that we want to add a constant vector to each row of a matrix. We could do it like this: **(cont.)**

```
# Add the vector v to each row of the matrix x with an explicit loop
for i in range(4):
    y[i, :] = x[i, :] + v

# Now y is the following
# [[ 2  2  4]
# [ 5  5  7]
# [ 8  8 10]
# [11 11 13]]
print(y)
```

Broadcasting

- This works; however, when the matrix **x** is very large, computing an explicit loop in Python could be slow.
- Note that adding the vector **v** to each row of the matrix **x** is equivalent to forming a matrix **vv** by stacking multiple copies of **v** vertically, then performing elementwise summation of **x** and **vv**. We could implement this approach like this:

```
import numpy as np

# We will add the vector v to each row of the matrix x,
# storing the result in the matrix y
x = np.array([[1,2,3], [4,5,6], [7,8,9], [10, 11, 12]])
v = np.array([1, 0, 1])
vv = np.tile(v, (4, 1)) # Stack 4 copies of v on top of each other
print(vv)               # Prints "[[1 0 1]
                        #          [1 0 1]
                        #          [1 0 1]
                        #          [1 0 1]]"

y = x + vv # Add x and vv elementwise
print(y)   # Prints "[[ 2  2  4
            #          [ 5  5  7]
            #          [ 8  8 10]
            #          [11 11 13]]"
```



Broadcasting

- Numpy broadcasting allows us to perform this computation without actually creating multiple copies of **v**. Consider this version, using broadcasting:

```
import numpy as np

# We will add the vector v to each row of the matrix x,
# storing the result in the matrix y
x = np.array([[1,2,3], [4,5,6], [7,8,9], [10, 11, 12]])
v = np.array([1, 0, 1])
y = x + v # Add v to each row of x using broadcasting
print(y) # Prints "[[ 2  2  4]
          #           [ 5  5  7]
          #           [ 8  8 10]
          #           [11 11 13]]"
```

- The line **y = x + v** works even though **x** has shape **(4, 3)** and **v** has shape **(3,)** due to broadcasting; this line works as if **v** actually had shape **(4, 3)**, where each row was a copy of **v**, and the sum was performed elementwise.

Broadcasting

- Here are some applications of broadcasting:

```
import numpy as np

# Compute outer product of vectors
v = np.array([1,2,3]) # v has shape (3,)
w = np.array([4,5])   # w has shape (2,)
# To compute an outer product, we first reshape v to be a column
# vector of shape (3, 1); we can then broadcast it against w to yield
# an output of shape (3, 2), which is the outer product of v and w:
# [[ 4  5]
#  [ 8 10]
#  [12 15]]
print(np.reshape(v, (3, 1)) * w)

# Add a vector to each row of a matrix
x = np.array([[1,2,3], [4,5,6]])
# x has shape (2, 3) and v has shape (3,) so they broadcast to (2, 3),
# giving the following matrix:
# [[2 4 6]
#  [5 7 9]]
print(x + v)
```

Broadcasting

- Here are some applications of broadcasting (**cont.**):

```
# Add a vector to each column of a matrix
# x has shape (2, 3) and w has shape (2,).
# If we transpose x then it has shape (3, 2) and can be broadcast
# against w to yield a result of shape (3, 2); transposing this result
# yields the final result of shape (2, 3) which is the matrix x with
# the vector w added to each column. Gives the following matrix:
# [[ 5  6  7]
#  [ 9 10 11]]
print((x.T + w).T)

# Another solution is to reshape w to be a column vector of shape (2, 1);
# we can then broadcast it directly against x to produce the same
# output.
print(x + np.reshape(w, (2, 1)))

# Multiply a matrix by a constant:
# x has shape (2, 3). Numpy treats scalars as arrays of shape ();
# these can be broadcast together to shape (2, 3), producing the
# following array:
# [[ 2  4  6]
#  [ 8 10 12]]
print(x * 2)
```

+ Numpy Documentation

- This brief overview has touched on many of the important things that you need to know about numpy, but is far from complete. Check out the [numpy reference](#) to find out much more about numpy.



[**https://numpy.org/doc/**](https://numpy.org/doc/)
[**https://docs.scipy.org/doc/numpy/reference/**](https://docs.scipy.org/doc/numpy/reference/)



Any Questions?